

THEORY OF DISTRIBUTED COMPUTING

Spanning Tree Construction
Computation in trees

A.A. 2023/24


Main References

"DESIGN AND ANALYSIS OF DISTRIBUTED ALGORITHMS"

Nicola Santoro
Wiley 2007 (available at the library)

Original slides by Paola Flocchini,
SITE, University of Ottawa, Canada
(rearranged by Manuela Montangero)

SLIDES CAN NOT BE REDISTRIBUTED




Spanning Tree Construction

A **spanning tree** T of a graph $G = (V, E)$ is an acyclic subgraph of G such that $T = (V, E')$ and $E' \subseteq E$.

Restrictions:

- Single initiator
- Bidirectional links
- Total reliability
- G connected




Spanning Tree: Protocol SHOUT

IDEAS?


Initially:


$$\forall x, \text{Tree-neighbors}(x) = \{ \}$$



At the end:

$$\forall x, \text{Tree-neighbors}(x) = \{ \text{neighbours in the spanning tree} \}$$





spanning
tree edge

Observation


At the end entities do not know the entire spanning tree, but only those edges that connect them to the neighbours in the spanning tree

Protocol SHOUT

1.




initiator




Q?
"Do you want to be
my neighbour
in the spanning tree ?"

2.




any other entity

If it is the first question:



If has already answered yes before:



Protocol SHOUT

State $S = \{\text{INITIATOR}, \text{IDLE}, \text{ACTIVE}, \text{DONE}\}$
Sinit = $\{\text{INITIATOR}, \text{IDLE}\}$ (possible initial states)
Sterm = $\{\text{DONE}\}$ (termination state)

Protocol per agent x

INITIATOR

spontaneously

```
root = true
Tree-neighbours(x) = {}
send Q to N(x)
counter = 0
become (ACTIVE)
```

Counts the number
of answers

IDLE

receiving(Q)

```
root = false
parent = sender
Tree-neighbours(x) = {sender}
send YES to parent
counter = 1
if counter = |N(x)|
  then
    become (DONE)
  else
    send Q to N(x) - {sender}
    become (ACTIVE)
```

Protocol SHOUT

State $S = \{\text{INITIATOR}, \text{IDLE}, \text{ACTIVE}, \text{DONE}\}$
Sinit = $\{\text{INITIATOR}, \text{IDLE}\}$ (possible initial states)
Sterm = $\{\text{DONE}\}$ (termination state)

ACTIVE

receiving(Q)

send NO **to sender**

receiving(YES)

Tree-neighbours(x) =

Tree-neighbours(x) \cup { **sender** }

counter = counter + 1

if counter = $|N(x)|$

then

become (DONE)

receiving(NO)

counter = counter + 1

if counter = $|N(x)|$


then

become (DONE)

For any other
pair
(state,event)
the
corresponding
action is nil

SHOUT: correctness and termination

- If x is in Tree-neighbours of y ,
then y is in Tree-neighbours of x
- If x sends YES to y ,
then y is in Tree-neighbours of x
and
is connected to the initiator by a chain of YES
- Every x (except the initiator) sends exactly one YES




The spanning graph defined by the Tree-neighbours relation is a connected tree containing all entities

Note: local termination

SHOUT: message complexity

Observation

SHOUT = FLOODING + REPLY




we expect

$$Message(SHOUT) = 2Message(FLOODING)$$

SHOUT: message complexity

Possible situations




Impossible situations




SHOUT: message complexity - worst case

Total number of Q



$(n-1)$

only one Q on the ST links



$2[m - (n-1)]$

on the other links

$$\text{Total} = (n-1) + 2[m - (n-1)] = 2m - n + 1$$

Total number of NO




$2[m - (n-1)]$

as many as Q---Q

SHOUT: message complexity - worst case

Total number of YES



exactly one yes only on the ST links

Total number of messages (Q + NO + YES)

$$\begin{aligned} \text{Message}(\text{SHOUT}) &= 2m - n + 1 + 2[m - (n-1)] + (n-1) \\ &= 4m - 2n + 2 \\ &= 2(2m - n + 1) = 2\text{Message}(\text{FLOODING}) \end{aligned}$$

Is it possible to reduce the number of messages?

Protocol SHOUT+: spanning tree construction without NO

State $S = \{\text{INITIATOR}, \text{IDLE}, \text{ACTIVE}, \text{DONE}\}$
Sinit = $\{\text{INITIATOR}, \text{IDLE}\}$ (possible initial states)
Sterm = $\{\text{DONE}\}$ (termination state)

Actions in reaction to events when in states

INITIATOR and IDLE

do not change

Protocol SHOUT+:

spanning tree construction without NO

State $S = \{\text{INITIATOR}, \text{IDLE}, \text{ACTIVE}, \text{DONE}\}$
Sinit = $\{\text{INITIATOR}, \text{IDLE}\}$ (possible initial states)
Sterm = $\{\text{DONE}\}$ (termination state)

ACTIVE

receiving(Q) ← to be interpreted as no

counter = counter + 1

if counter = $|N(x)|$

then

become (DONE)

receiving(YES)

Tree-neighbours(x) =

Tree-neighbours(x) $\cup \{\text{sender}\}$

counter = counter + 1



if counter = $|N(x)|$

then

become (DONE)

SHOUT+: message complexity

On each link there will be exactly two messages



$$Message(SHOUT+) = 2m$$

much better than

$$2(2m - n + 1) = Message(SHOUT)$$

Spanning Tree Construction by Traversal

A spanning tree can be build
by traversing the graph

TRAVERSAL PROBLEM

Initially all entities are in the same
unvisited state
except for one that is *visited* and is the *initiator*
The goal is to make all entities
visited sequentially (one at the time)

Traversal protocol

Distributed algorithm that, starting from the
initiator, uses a special message (*token*) and reaches
every entity *sequentially*. Once an entity receives
the token is considered visited.

A depth-first traversal of a graph builds a spanning tree of the graph

Spanning Tree Construction by Traversal

DEPTH-FIRST traversal

“The graph is traversed trying to forward (the token) as long as possible”

Restrictions

- Single initiator
- Bidirectional links
- Connectivity
- Total reliability

Spanning Tree Construction by Traversal

Traversal protocol

● Visited node


🌀 Forward Token

🌀 Back-edge Token

🌀 Return Token

1. When receiving the **Forward Token** the first time:

- remember who sent the token
- send the **Forward Token** to ONE unvisited neighbour
- wait for token to return




2. When receiving the token again:

if is **ReturnToken**:

if there still are unvisited neighbours with no back-edge:

- send the **Forward Token** to ONE unvisited neighbour
- wait for token to return

otherwise



Spanning Tree Construction by Traversal

Traversal protocol

● Visited node


🌀 Forward Token

🌀 Back-edge Token

🌀 Return Token

1. When receiving the **Forward Token** the first time:

- remember who sent the token
- send the **Forward Token** to ONE unvisited neighbour
- wait for token to return



2. When receiving the token again:


if is **ReturnToken**:

if there still are unvisited neighbours with no back-edge:

- send the **Forward Token** to ONE unvisited neighbour
- wait for token to return

otherwise

send **ReturnToken** to the one from which it first received the **Forward Token**



Spanning Tree Construction by Traversal

Traversal protocol

● Visited node


⦿ Forward Token

⦿ Back-edge Token

⦿ Return Token

1. When receiving the **Forward Token** the first time:

- remember who sent the token
- send the **Forward Token** to ONE unvisited neighbour
- wait for token to return



2. When receiving the token again:

if is **Forward Token**:

send back to sender **Back-edge Token** and classify the link as **back-edge**




Spanning Tree Construction by Traversal

Traversal protocol



1. When receiving the **Forward Token** the first time:

- remember who sent the token
- send the **Forward Token** to ONE unvisited neighbour
- wait for token to return



2. When receiving the token again:

if is **Forward Token**:


send back to sender **Back-edge Token** and classify the link as **back-edge**

if is **Back-edge Token**:







proceed as with **Return Token**




Depth-First Traversal Example



Depth-First Traversal Example



Spanning Tree Construction by Traversal



Removing back-edges
we have a spanning tree

Who is the root?

The root of the tree is the initiator

Who is entity x's
parent?

The parent of entity x is the one from
which it first received the token

Who are entity x's
children?

The children of entity x are the neighbours
that are not connected by a back-edge

Spanning Tree Construction by Traversal

State $S = \{\text{INITIATOR}, \text{IDLE}, \text{VISITED}, \text{DONE}\}$
Sinit = $\{\text{INITIATOR}, \text{IDLE}\}$ (possible initial states)
Sterm = $\{\text{DONE}\}$ (termination state)

Protocol per agent x

INITIATOR
spontaneously
Unvisited := $N(x)$
initiator := true
VISIT

Picks one element in
Unvisited and
eliminates the element
from the set

```
procedure VISIT
  if |Unvisited| > 0
  then
    next := pick(Unvisited)
    send Token to next
    become (VISITED)
  else
    if not(initiator)
    then
      send ReturnToken to entry
    become (DONE)
```

Spanning Tree Construction by Traversal

State $S = \{\text{INITIATOR}, \text{IDLE}, \text{VISITED}, \text{DONE}\}$
Sinit = $\{\text{INITIATOR}, \text{IDLE}\}$ (possible initial states)
Sterm = $\{\text{DONE}\}$ (termination state)

Protocol per agent x

IDLE

receiving(Token)

entry := **sender**

Unvisited := $N(x) \setminus \{\text{sender}\}$

initiator := false

VISIT

procedure **VISIT**

if |Unvisited| > 0

then

next := pick(Unvisited)

send Token **to** next

become (VISITED)

else

if not(initiator)

then

send ReturnToken **to** entry

become (DONE)

Picks one element in
Unvisited and
eliminates the element
from the set

Spanning Tree Construction by Traversal

State $S = \{\text{INITIATOR}, \text{IDLE}, \text{VISITED}, \text{DONE}\}$
Sinit = $\{\text{INITIATOR}, \text{IDLE}\}$ (possible initial states)
Sterm = $\{\text{DONE}\}$ (termination state)

Protocol per agent x

VISITED

receiving(Token)

Unvisited := Unvisited \ {**sender**}

send BackEdgeToken **to sender**

receiving(ReturnToken)

VISIT

receiving(BackEdgeToken)

VISIT

a descendant in the spanning tree is also a neighbour of x .
 x will not visit it later on.

a neighbour terminated its visit,
 x continues with next neighbour


the neighbour has already received the token before,
 x continues with next neighbour

When **DONE**, agent x will
not receive messages any more


Spanning Tree Construction by Traversal

How many messages to perform depth-first traversal?


Given a link



followed by



XOR



followed by



MESSAGE COMPLEXITY = $2m$

Is it possible to reduce message complexity?

No, $Message(DFT(G)) \in \Omega(m)$

(the proof is analogous to the one given for broadcast)

Spanning Tree Construction by Traversal

How much time to perform depth-first traversal?

Since traversal is sequential, and $2m$ messages are sent sequentially...

$$\text{TIME COMPLEXITY} = 2m$$

Can we do better?

Time complexity lower bound
 $\text{Time}(DFT(G)) \geq n - 1$

each node has to be visited sequentially

Spanning Tree Construction by Traversal

Improving time complexity

Each entity MUST receive the token at least once
BUT
in our protocol each entity receives the token once
from each neighbour

IDEA

Avoid sending the token on back-edges

Spanning Tree Construction by Traversal

Improving time complexity




Avoid sending the token on back-edges

At any time, each entity
has a set of visited
neighbors
(initially empty)

... send "Visited" msgs
to non visited neighbors and...

... wait for "Acks" msgs...

When receiving the token
the first time...*



* holds for initiator as well


... then proceed as before considering only non visited neighbors.

Spanning Tree Construction by Traversal

Improving time complexity


Avoid sending the token on back-edges

When receiving a "Visited" message...



Visited →

... send "Ack" msg to sender
and eliminate sender from
the set of unvisited neighbors...



← Ack


... then proceed as before, always considering only non visited neighbors.

Spanning Tree Construction by Traversal


Improving time complexity

Avoid sending the token on back-edges

Messages through links



OR



Sequential:

$2(n-1)$ message chain long

Concurrent to/from neighbors:

2 message chain long for each neigh.

longest message chain

$$2n - 2 + 2n = 4n - 2$$


TIME COMPLEXITY $\in O(n)$

Spanning Tree Construction by Traversal


Avoid sending the token on back-edges

message complexity?

Messages through links



OR



For each entity (except initiator):
2 messages with token per link
in the tree

For each link: 2 messages
("visited" will not be sent twice on link)

number of messages
 $2(n-1) + 2m$

MESSAGE COMPLEXITY $\in O(m)$

Spanning Tree Construction

Message(SHOUT+) = $2m$

Which one?

Message(DFT) = $2m (+ 2(n-1))$

Different techniques construct different spanning tree

The same protocol on the same graph might produce different spanning trees when executed at different times


Using SHOUT, it is impossible to predict which spanning tree will be constructed

In general DFT constructs a tree with terrible diameter


Spanning Tree Construction

Ideally it is desirable to have a spanning tree with SMALL diameter

WHY?



One possible
spanning tree T
 $D(T) = 4$



Another possible
spanning tree T'
 $D(T') = 2$

Which one is more desirable for broadcasting?

Spanning Tree Construction

Ideally it is desirable to have a spanning tree
with SMALL diameter

HOW?





Broadcast-Tree Construction:

1. Determine a *center of G*
2. Construct a **breadth-first spanning**
tree rooted in the center

1 and 2 are both expensive tasks

We will not go into further details

What happens with SHOUT if there are multiple initiators ?




ST edge
non ST edge

Protocol SHOUT produces a forest

What happens if there are multiple initiators ?

In general an entity does not know if there are other initiators



- Devise a different protocol

impossible if deterministic
and entities do not have
unique identifiers

OR


- ELECT an entity (LEADER) to be the unique initiator

Spanning Tree with multiple initiators

Additional restriction: UNIQUE IDs

IDEA 1: Multiple Spanning Tree

Each initiator constructs its own spanning tree with a single-initiator protocol and uses the IDs of the initiators to distinguish between the different constructions



Message cost depends on the number of initiators and used protocol.
In general is expensive.

Spanning Tree with multiple initiators

Additional restriction: UNIQUE IDs

IDEA 2: Selective construction

Each initiator starts the construction of its own spanning tree with a single-initiator protocol using their IDs to identify their spanning tree. Entities will eventually stop working for all but one constructions, keeping the spanning tree of the initiator with smaller ID.



Entities might re-execute the protocol several times


Need of a termination notification

Before talking about leader election...

Computation in Trees


Entities are aware of belonging to a tree network

Acyclic graph
 n entities
 $n-1$ edges



Rooted Tree

Sense of direction: up-down



Unrooted Tree

Computation in Unrooted Trees

Restrictions

Bidirectional links

Connectivity

FIFO messages

Full reliability

Knowledge of the topology

Each entity knows the network is a tree


Each entity knows if

it is a leaf:



Only one
neighbour

or an internal node:



Computation in Unrooted Trees

SATURATION TECHNIQUE

EXAMPLE: Minimum Finding


Every entity x has an input value $v(x)$
(not necessarily distinct)

At the end each entity must know if its
value is the smallest or not

Computation in Unrooted Trees

SATURATION TECHNIQUE

EXAMPLE: Minimum Finding




Leaves start the computation sending their value to their unique neighbor

Computation in Unrooted Trees

SATURATION TECHNIQUE

EXAMPLE: Minimum Finding




Internal entities wait for all but one neighbor to
send a message,
then compute minimum and send to last neighbor

Computation in Unrooted Trees

SATURATION TECHNIQUE

EXAMPLE: Minimum Finding




Internal entities wait for all but one neighbor to
send a message,
then compute minimum and send to last neighbor

Computation in Unrooted Trees

SATURATION TECHNIQUE

EXAMPLE: Minimum Finding




Two entities receive a message from all neighbors
and send the minimum back

Computation in Unrooted Trees

SATURATION TECHNIQUE

EXAMPLE: Minimum Finding




All internal entities send the message containing the minimum to the neighbors it first received messages from

Computation in Unrooted Trees

SATURATION TECHNIQUE

EXAMPLE: Minimum Finding




All internal entities send the message containing the minimum to the neighbors it first received messages from

Computation in Unrooted Trees

SATURATION TECHNIQUE

EXAMPLE: Minimum Finding



At the end, all entities know the minimum value and can decide (if they hold it or not)

Computation in Unrooted Trees

SATURATION TECHNIQUE

Full Saturation

Can be autonomously and independently started by any number of initiators

Activation stage:

Started by all initiators: all entities are activated

Saturation stage:

Started by leaves

At the end, one pair of neighbor entities is selected

Resolution stage:

Started by selected (saturated) entities

Computation in Unrooted Trees


SATURATION TECHNIQUE

Activation stage:

Started by all initiators: all entities are activated

Wake-up started by initiators

Within a finite time all entities
become **active**



Computation in Unrooted Trees



SATURATION TECHNIQUE

Saturation stage:



Started by leaves

At the end, one pair of neighbor entities is selected

Leaves send a saturation message to their only neighbor



Internal entities wait $|N(.)|-1$ saturation messages and then send a saturation message to the remaining neighbor



Computation in Unrooted Trees



SATURATION TECHNIQUE

Saturation stage:

Started by leaves

At the end, one pair of neighbor entities is selected

If a **processing** entity receives a saturation message
it becomes **saturated**



Computation in Unrooted Trees


SATURATION TECHNIQUE

Resolution stage:

Started by selected (saturated) entities

Depends on the application

Usually is a **notification**




Computation in Unrooted Trees

SATURATION TECHNIQUE

LEMMA

Exactly two processing entities become saturated, and they are neighbors



Different execution might result into a different pair of saturated entities, depending on communication delays

Computation in Unrooted Trees


SATURATION TECHNIQUE

MESSAGE COMPLEXITY

Activation:

Saturation:

Resolution:



Computation in Unrooted Trees


SATURATION TECHNIQUE

MESSAGE COMPLEXITY

Activation: Worst Case
n initiators

Saturation:

Resolution:



$2 (n-1)$ messages


Computation in Unrooted Trees

SATURATION TECHNIQUE

MESSAGE COMPLEXITY

Activation: $\leq 2(n-1)$

Saturation:



Resolution:

$(n-1) + 1 = n$ messages

Computation in Unrooted Trees


SATURATION TECHNIQUE

MESSAGE COMPLEXITY

Activation: $\leq 2(n-1)$

Saturation: n

Resolution:



$n-2$ messages

Computation in Unrooted Trees


SATURATION TECHNIQUE

MESSAGE COMPLEXITY

Activation: $\leq 2(n-1)$

Saturation: n

Resolution: $n-2$



$$\text{Total} \leq 2(n-1) + n + n - 2 = 4n - 4$$

Computation in Unrooted Trees

SATURATION TECHNIQUE


Saturation can be used to solve a wide set of problems

Distributed Function Evaluation

compute a function whose arguments
are initially distributed among entities

- Minimum Finding
- Cardinal statistics
- Find eccentricity
- Center finding
- Finding a median
- Finding diametral path

Computation in Rooted Trees



Rooted Tree

Theorem

Without unique ID
it is impossible to
root an unrooted tree

Sense of direction
up-down

Each entity knows who are the children
and who is the parent

There is a natural leader,
the root

Protocols are started by
the root with a broadcast

“Saturation” of the root is
achieved by convergecast

Convergecast:

1. Leaves send msg to parent
2. internal entities send to parent after receiving from all children